

Hindi Vidya Prachar Samiti's
RAMNIRANJAN JHUNJHUNWALA COLLEGE
OF ARTS, SCIENCE & COMMERCE
(EMPOWERED AUTONOMOUS)

Big Data Analytics



Name: Siddhi Shashikant Chavan

Roll No.: 711

Class: MSc Data Science and Artificial Intelligence Part-I



Hindi Vidya Prachar Samiti's Ramniranjan Jhunjhunwala College of Arts, Science and Commerce

Department of Data Science and Artificial Intelligence

CERTIFICATE

This is to certify Ms. Siddhi Shashikant Chavan of Msc. Data Science and Artificial Intelligence Roll No 711 has successfully completed the practical of Big Data Analytics during the Academic Year 2023-2024.

Date :

(Prof. Mujtaba Shaikh)

Prof-In-Charge

External Examiner

INDEX

Sr. No.	Practical Name	Date	Remark
1	Perform ETL on student data and store transformed information in Data Warehouse.	30 Nov 2023	
2	Perform basic commands on Hadoop distributed file system.	16 Dec 2023	
3	Implementation of MapReduce program to find word counts.	04 Jan 2024	
4	Perform linear regression on the given dataset using PySpark on Databricks.	13 Jan 2024	
5	Perform customer Churn analysis on the given dataset using ML Algorithm in PySpark.	19 Jan 2024	
6	Implementation of Cluster analysis using K-Means Clustering Algorithm & also find the optimal number of cluster using elbow method.	27 Jan 2024	
7	Demonstrate the working of internal tables & external tables in HIVE.	02 Feb 2024	
8	Perform Data Ingestion using Apache Sqoop tool: a) MySQL to HDFS (import) b) MySQL to HIVE (import) c) MySQL to HBASE (import) d) HDFS to MySQL (export) e) HIVE to MySQL (export)	15 Feb 2024	
9	Perform data processing using pig latin.	29 Feb 2024	

Practical 1

Perform ETL on student data and store transformed information in Data Warehouse.

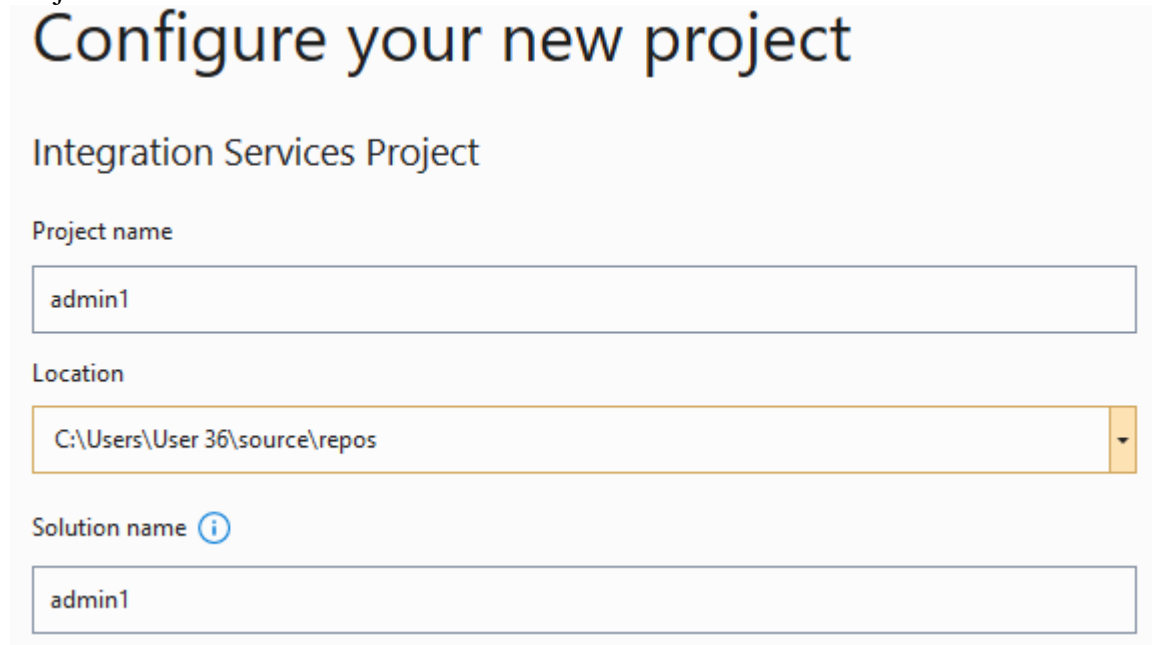
Create dataset

Use Visual Studio 2019

Create a new project

Integration Services project

Project name



Configure your new project

Integration Services Project

Project name

admin1

Location

C:\Users\User 36\source\repos

Solution name ⓘ

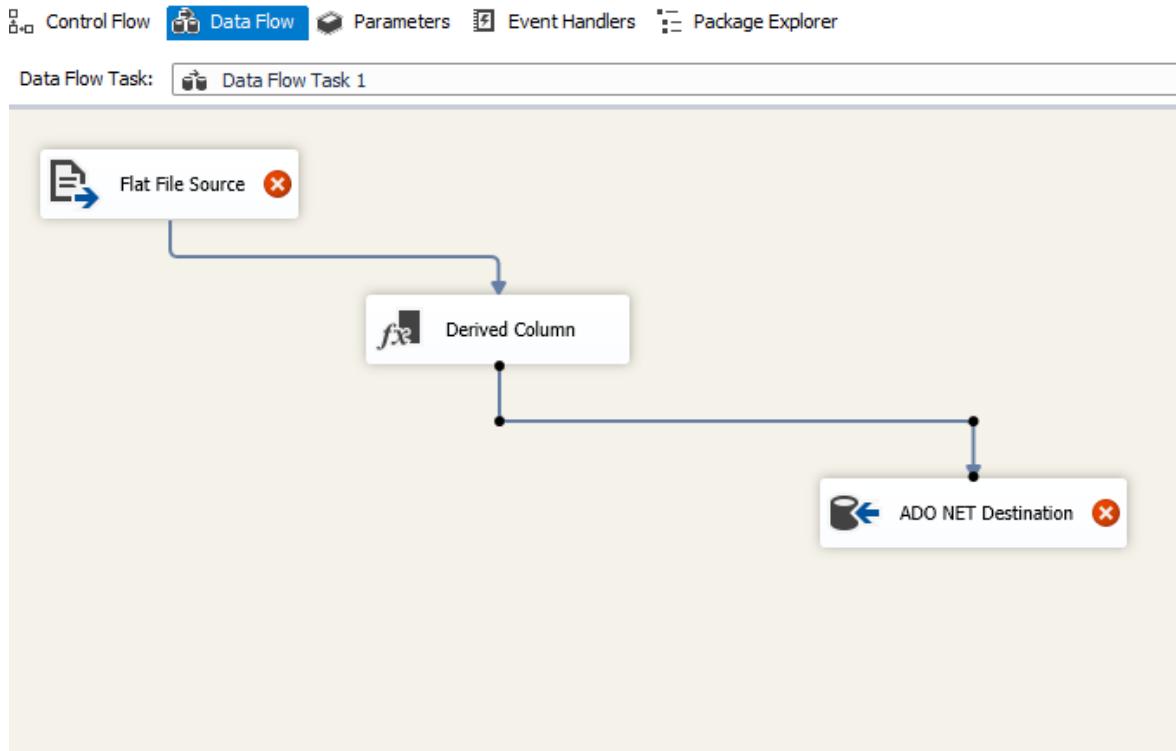
admin1

Data Flow Task

Flat File

Derived Column

ADO NET Destination



Flat File source

-Connection manager-New-Browser-Preview Derived Column

-Derived Column Name-Derived Column-expression

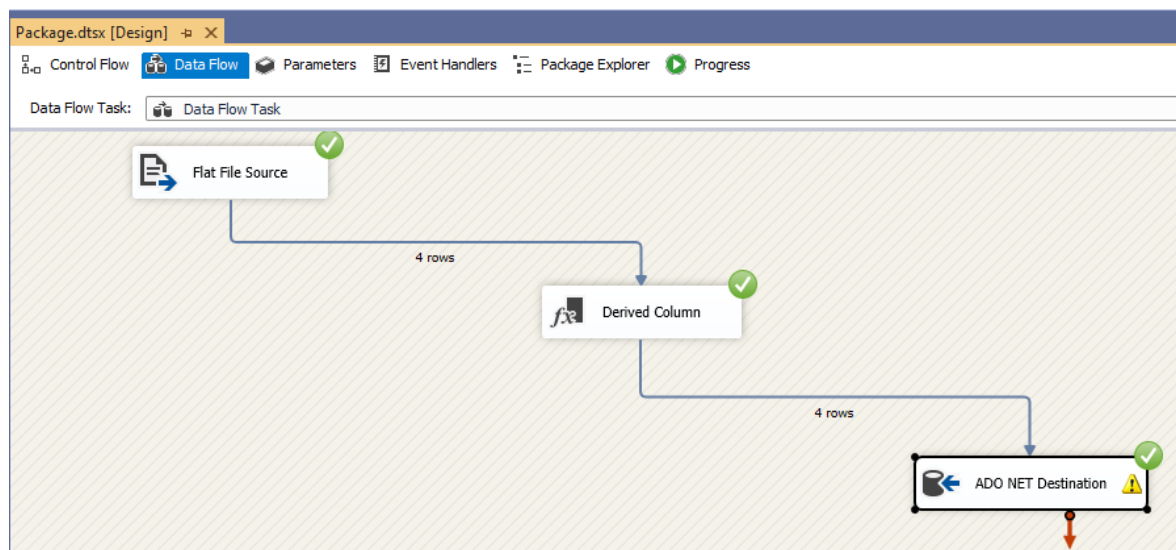
Microsoft SQL Server Management Studio

Create Database

Create Table

ADO NET Destination

-Connection Manager-New



Select Top 1000 Rows

SQLQuery1.sql - DE...RMGJ\User 36 (57) - X

/****** Script for SelectTopNRows command from SSMS *****/

```
SELECT TOP (1000) [name]
      ,[id]
      ,[age]
      ,[gender]
FROM [admin].[dbo].[Table_1]
```

	name	id	age	gender
1	ANU	1	12	f
2	SITA	2	45	f
3	GEETA	3	76	f
4	RAM	4	56	m
5	ANU	1	12	f
6	SITA	2	45	f
7	GEETA	3	76	f
8	RAM	4	56	m

Practical 2

Perform basic commands on Hadoop distributed file system.

```
[cloudera@quickstart ~]$ sudo su
[root@quickstart cloudera]# jps
```

```
8130
5502 SecondaryNameNode
5676 JobHistoryServer
5286 JournalNode
6915 RunJar
7807 Bootstrap
5792 NodeManager
6634 ThriftServer
5118 QuorumPeerMain
8170
7458 HRegionServer
5190 DataNode
8104 Bootstrap
6041 ResourceManager
5626 Bootstrap
6315 HMaster
6452 RESTServer
7335 HistoryServer
9424 Jps
6757 RunJar
7306 Bootstrap
5371 NameNode
```

```
[root@quickstart cloudera]# su cloudera
[cloudera@quickstart ~]$ hadoop version
```

```
Hadoop 2.6.0-cdh5.13.0
Subversion http://github.com/cloudera/hadoop -r 42e8860b182e55321bd5
Compiled by jenkins on 2017-10-04T18:08Z
Compiled with protoc 2.5.0
From source with checksum 5e84c185f8a22158e2b0e4b8f85311
This command was run using /usr/lib/hadoop/hadoop-common-2.6.0-cdh5.
```

```
[cloudera@quickstart ~]$ hdfs fsck /
```

```
Connecting to namenode via http://quickstart.c
FSCK started by cloudera (auth:SIMPLE) from /f
.....
.....
.....
.....
.....
.....
.....
.....
.....
Status: HEALTHY
Total size:      861286254 B (Total open files
Total dirs:     79
Total files:    931
Total symlinks:                               0 (Files curre
Total blocks (validated):                     929 (avg. bloc
Minimally replicated blocks:                  929 (100.0 %)
Over-replicated blocks:                       0 (0.0 %)
Under-replicated blocks:                      0 (0.0 %)
Mis-replicated blocks:                        0 (0.0 %)
Default replication factor:                   1
Average block replication:                    1.0
Corrupt blocks:                              0
```

```
[cloudera@quickstart ~]$ hdfs dfs -mkdir /bigdata
[cloudera@quickstart ~]$ hdfs dfs -touchz /bigdata/mydata.dat
[cloudera@quickstart ~]$ hdfs dfs -ls /bigdata/
```

```
Found 1 items
-rw-r--r--  1 cloudera supergroup          0 2024-02-14 23:06 /bigdata/mydata.dat
```

```
[cloudera@quickstart ~]$ touch linux.txt
```

```
[cloudera@quickstart ~]$ ls
cloudera-manager  data  Documents  eclipse  express-deployment.json  lib  Music
cm_api.py         Desktop  Downloads  enterprise-deployment.json  kerberos  linux.txt  parcels
```

```
[cloudera@quickstart ~]$ hdfs dfs -put linux.txt /
```

```
[cloudera@quickstart ~]$ hdfs dfs -ls /
```

```
Found 8 items
drwxrwxrwx  - hdfs      supergroup          0 2017-10-23 09:15 /bench
drwxr-xr-x  - cloudera supergroup          0 2024-02-14 23:06 /bigdata
drwxr-xr-x  - hbase    supergroup          0 2024-02-14 22:19 /hbase
-rw-r--r--  1 cloudera supergroup          0 2024-02-14 23:36 /linux.txt
drwxr-xr-x  - solr     solr              0 2017-10-23 09:18 /solr
drwxrwxrwt  - hdfs     supergroup          0 2024-02-14 22:19 /tmp
drwxr-xr-x  - hdfs     supergroup          0 2017-10-23 09:17 /user
drwxr-xr-x  - hdfs     supergroup          0 2017-10-23 09:17 /var
```

```
[cloudera@quickstart ~]$ hdfs dfs -put linux.txt /bigdata
```

```
[cloudera@quickstart ~]$ hdfs dfs -ls /bigdata
```

```
Found 8 items
drwxrwxrwx  - hdfs      supergroup          0 2017-10-23 09:15 /bench
drwxr-xr-x  - cloudera supergroup          0 2024-02-14 23:38 /bigdata
drwxr-xr-x  - hbase    supergroup          0 2024-02-14 22:19 /hbase
-rw-r--r--  1 cloudera supergroup          0 2024-02-14 23:36 /linux.txt
drwxr-xr-x  - solr     solr              0 2017-10-23 09:18 /solr
drwxrwxrwt  - hdfs     supergroup          0 2024-02-14 22:19 /tmp
drwxr-xr-x  - hdfs     supergroup          0 2017-10-23 09:17 /user
drwxr-xr-x  - hdfs     supergroup          0 2017-10-23 09:17 /var
```

```
[cloudera@quickstart ~]$ hdfs dfs -appendToFile - /bigdata/data.bat
```

```
twtgrvfvsferty
jyyhyhjjk
```

Type Control D

```
[cloudera@quickstart ~]$ hdfs dfs -cat /bigdata/data.bat
```

```
twtgrvfvsferty
jyyhyhjjk
```


Practical 3

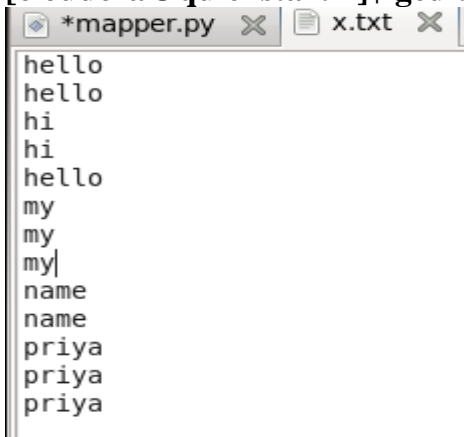
Implementation of MapReduce program to find word counts.

[cloudera@quickstart ~]\$ gedit mapper.py

```
#!/usr/bin/python

import sys
for line in sys.stdin:
    line = line.strip()
    hello = line.split()
    for x in hello:
        print '%s\t%s'%(x,1)
```

[cloudera@quickstart ~]\$ gedit x.txt



```
hello
hello
hi
hi
hello
my
my
my
name
name
priya
priya
priya
```

[cloudera@quickstart ~]\$ cat x.txt | python mapper.py

```
hello 1
hello 1
hi 1
hi 1
hello 1
my 1
my 1
my 1
name 1
name 1
priya 1
priya 1
priya 1
```

[cloudera@quickstart ~]\$ gedit reducer.py

```
#!/usr/bin/python
import sys
from operator import itemgetter
current_word=None
current_count=0
word=None
for line in sys.stdin:
    line = line.strip()
    word,count=line.split('\t',1)
    try:
        count=int(count)
    except ValueError:
        continue
    if current_word==word:
        current_count+=count
    else:
        if current_word:
            print '%s\t%s'%(current_word,current_count)
            current_count=count
            current_word=word
if current_word==word:
    print '%s\t%s'%(current_word,current_count)
```

cloudera@quickstart ~]\$ cat x.txt | python mapper.py| sort | python reducer.py

```
hello 3
hi 2
my 3
name 2
priya 3
cloudera@qu
```

Practical 4

Perform linear regression on the given dataset using PySpark on Databricks.

```
from pyspark.sql import SparkSession
ss=SparkSession.builder.getOrCreate()
df=ss.read.csv('/FileStore/tables/housesales.csv',inferSchema=True,header=True)
df.show()
```

```
+---+-----+-----+-----+
| No|area-sq|on_of_bed| price|
+---+-----+-----+-----+
|  1|    500|         2|105000|
|  2|    550|         2|120000|
|  3|    440|         1| 96000|
|  4|    300|         1| 90000|
|  5|    450|         1|980000|
|  6|    600|         2|125000|
|  7|    700|         4|150000|
|  8|    650|         3|140000|
|  9|    510|         2|110000|
+---+-----+-----+-----+
```

```
df.printSchema()
```

```
root
 |-- No: integer (nullable = true)
 |-- area-sq: integer (nullable = true)
 |-- on_of_bed: integer (nullable = true)
 |-- price: integer (nullable = true)
```

```
from pyspark.ml.feature import VectorAssembler
```

```
#fa =feature assemble
```

```
fa=VectorAssembler(inputCols=["area-
sq","on_of_bed"],outputCol="Indep_feature")
```

```
output=fa.transform(df)
```

```
output.show()
```

► (1) Spark Jobs

No	area-sq	on_of_bed	price	Indep_feature
1	500	2	105000	[500.0,2.0]
2	550	2	120000	[550.0,2.0]
3	440	1	96000	[440.0,1.0]
4	300	1	90000	[300.0,1.0]
5	450	1	980000	[450.0,1.0]
6	600	2	125000	[600.0,2.0]
7	700	4	150000	[700.0,4.0]
8	650	3	140000	[650.0,3.0]
9	510	2	110000	[510.0,2.0]

```
find_data=output.select("indep_feature","price")
```

```
find_data.show()
```

► (1) Spark Jobs

indep_feature	price
[500.0,2.0]	105000
[550.0,2.0]	120000
[440.0,1.0]	96000
[300.0,1.0]	90000
[450.0,1.0]	980000
[600.0,2.0]	125000
[700.0,4.0]	150000
[650.0,3.0]	140000
[510.0,2.0]	110000

```
from pyspark.ml.regression import LinearRegression
```

```
train_data,test_data=find_data.randomSplit([0.75,0.25])
```

```
reg=LinearRegression(featuresCol="indep_feature",labelCol="price")
```

```
reg=reg.fit(train_data)
```

```
reg.coefficients
```

```
reg.intercept
```

```
pred_result=reg.evaluate(test_data)
```

```
pred_result.predictions.show()
```

indep_feature	price	prediction
[450.0,1.0]	980000	98435.77981651352
[700.0,4.0]	150000	151348.62385321147

```
pred_result.meanAbsoluteError,pred_result.meanSquaredError
```

Practical 5

Perform customer Churn analysis on the given dataset using ML Algorithm in PySpark.

```
from pyspark.sql import SparkSession
ss=SparkSession.builder.getOrCreate()
ss
df=ss.read.csv(
"/FileStore/tables/WA_Fn_UseC__Telco_Customer_Churn_csv__WA_Fn_UseC_
_Telco_Customer_Churn-1.csv",inferSchema=True,header=True)
df.display()
```

Table ▾ +

	customerID ▲	gender ▲	SeniorCitizen ▲	Partner ▲	Dependents ▲	tenure ▲
1	7590-VHVEG	Female	0	Yes	No	1
2	5575-GNVDE	Male	0	No	No	34
3	3668-QPYBK	Male	0	No	No	2
4	7795-CFOCW	Male	0	No	No	45
5	9237-HQITU	Female	0	No	No	2
6	9305-CDSKC	Female	0	No	No	8
7	1452-KIOVK	Male	0	No	Yes	22

```
df.printSchema()
df=df.drop('customerID')
df.display()
```

Table ▾ +

	gender ▲	SeniorCitizen ▲	Partner ▲	Dependents ▲	tenure
1	Female	0	Yes	No	1
2	Male	0	No	No	34
3	Male	0	No	No	2
4	Male	0	No	No	45
5	Female	0	No	No	2
6	Female	0	No	No	8
7	Male	0	No	Yes	22

```
from pyspark.ml.feature import StringIndexer
gen=StringIndexer(inputCol='gender',outputCol='new_gender')
```

```
df.show()
```

```
ce|No internet service|      Two year|      No|Ci
| Male|              0|      Yes|      No|  58|
es|              Yes|      One year|      No|Ci
| Male|              0|      No|      No|  49|
es|              Yes|Month-to-month|      Yes|Bi
| Male|              0|      No|      No|  25|
es|              Yes|Month-to-month|      Yes|
|Female|              0|      Yes|      Yes|  69|
es|              Yes|      Two year|      No|Ci
|Female|              0|      No|      No|  52|
ce|No internet service|      One year|      No|
| Male|              0|      No|      Yes|  71|
es|              Yes|      Two year|      No|Bi
|Female|              0|      Yes|      Yes|  10|
No|              No|Month-to-month|      No|Ci
|Female|              0|      No|      No|  21|
No|              Yes|Month-to-month|      Yes|
```

```
df.display()
```

Table ▾ +

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneServi
1	Female	0	Yes	No	1	No
2	Male	0	No	No	34	Yes
3	Male	0	No	No	2	Yes
4	Male	0	No	No	45	No
5	Female	0	No	No	2	Yes
6	Female	0	No	No	8	Yes
7	Male	0	No	Yes	22	Yes

↓ 7 043 rows | 1 24 seconds runtime

```
from pyspark.ml.feature import VectorAssembler
```

```
trans_col=trans.fit(df).transform(df)
```

```
trans_col.display()
```

Table ▾ +

	gender	SeniorCitizen	Partner	Dependents
1	Female	0	Yes	No
2	Male	0	No	No
3	Male	0	No	No
4	Male	0	No	No
5	Female	0	No	No
6	Female	0	No	No
7	Male	0	No	Yes


Practical 6

Implementation of Cluster analysis using K-Means Clustering Algorithm & also find the optimal number of cluster using elbow method.

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
df=pd.read_csv("student.csv")
Df
```

	cgpa	ML
0	5.13	88
1	5.90	113
2	8.36	93
3	8.27	97
4	5.45	110
...
195	4.68	89
196	8.57	118
197	5.85	112
198	6.23	108
199	8.82	117

```
plt.scatter(df[cgpa],df[ML] ,color="skyblue")
wss=[]
for i in range(1,11):
    km=KMeans(n_clusters=i)
    km.fit(df)
    wss.append(km.inertia_)
plt.plot(range(1,11),wss)
```



```
x=df.iloc[:,:].values
```

X

```
array([[ 5.13,  88.],
       [ 5.9,  93.],
       [ 8.36, 103.],
       [ 8.27,  97.],
       [ 5.45, 110.],
       [ 5.88, 109.],
       [ 8.41,  98.],
       [ 8.8, 115.],
       [ 5.79, 110.],
       [ 8.09,  94.],
       [ 4.6,  86.],
       [ 6.1, 106.],
       [ 8.16,  97.]])
```

```
km=KMeans(n_clusters=4)
km.fit(x)
```

```
KMeans(n_clusters=4)
```

```
y_pred=km.predict(x)
y_pred
```

```
array([[2, 1, 0, 0, 1, 1, 0, 3, 1, 0, 2, 1, 0, 2, 1, 0, 1, 0, 1, 1, 0, 2,
0, 2, 2, 0, 2, 3, 0, 1, 3, 1, 3, 1, 0, 0, 3, 1, 2, 1, 2, 0, 0, 2,
3, 3, 0, 1, 3, 1, 2, 2, 3, 0, 3, 1, 1, 3, 1, 3, 1, 0, 0, 3, 2, 3,
0, 2, 1, 0, 1, 3, 0, 2, 1, 3, 1, 3, 2, 0, 0, 3, 1, 2, 3, 2, 3, 1,
3, 1, 3, 3, 0, 2, 0, 0, 3, 0, 2, 3, 1, 2, 2, 3, 2, 2, 0, 2, 3, 3,
0, 3, 1, 1, 0, 3, 0, 1, 3, 2, 2, 1, 0, 3, 0, 2, 0, 1, 2, 0, 0, 1,
2, 2, 1, 3, 1, 2, 0, 0, 0, 2, 1, 2, 2, 3, 2, 3, 1, 2, 3, 2, 3, 3,
2, 0, 1, 3, 1, 0, 2, 3, 1, 0, 3, 2, 1, 2, 2, 3, 3, 1, 3, 2, 2, 0,
3, 1, 2, 3, 3, 1, 1, 1, 0, 2, 0, 0, 3, 1, 0, 0, 2, 2, 0, 2, 3, 1,
1, 3])
```

```
x[y_pred==0]
```

```
array([[ 8.36,  93. ],
[ 8.27,  97. ],
[ 8.41,  98. ],
[ 8.09,  94. ],
[ 8.16,  97. ],
[ 8.31,  95. ],
[ 7.87,  91. ],
[ 7.47,  98. ],
[ 7.78,  92. ],
[ 7.93,  98. ],
[ 8.04,  94. ],
[ 7.77,  96. ]])
```

```
X
```


```
array([[ 5.13,  88. ],
[ 5.9 , 113. ],
[ 8.36,  93. ],
[ 8.27,  97. ],
[ 5.45, 110. ],
[ 5.88, 109. ],
[ 8.41,  98. ],
[ 8.8 , 115. ],
[ 5.79, 110. ],
[ 8.09,  94. ],
[ 4.6 ,  86. ],
[ 6.1 , 110. ],
[ 8.16,  97. ],
[ 5.67, 100. ]])
```

```
x[y_pred==1]
```

```
array([[ 5.9 , 113. ],
[ 5.45, 110. ],
[ 5.88, 109. ],
[ 5.79, 110. ],
[ 6.1 , 110. ],
[ 5.71, 108. ],
[ 5.5 , 111. ],
[ 6.05, 111. ],
[ 5.84, 113. ],
[ 5.43, 106. ],
[ 6.01, 112. ],
[ 5.32, 106. ],
[ 5.91, 108. ],
[ 5.57, 113. ],
[ 6.4 , 108. ],
[ 5.67, 100. ]])
```

```
plt.scatter(x[y_pred==0,0],x[y_pred==0,1] ,color="red")
plt.scatter(x[y_pred==1,0],x[y_pred==1,1] ,color="blue")
plt.scatter(x[y_pred==2,0],x[y_pred==2,1] ,color="yellow")
plt.scatter(x[y_pred==3,0],x[y_pred==3,1] ,color="skyblue")
```

```
:matplotlib.collections.PathCollection at 0x212bf:
```



Practical 7

Demonstrate the working of internal and external table in Hive.

For first terminal

```
[cloudera@quickstart ~]$ hive home
```

```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
```

```
hive> show databases;
```

```
OK
```

```
default
```

```
Time taken: 0.463 seconds, Fetched: 1 row(s)
```

```
hive> create table student(roll_no int, name string,dept string);
```

```
hive> create database info;
```

```
hive> use info;
```

```
hive> create table student(roll_no int, name string,dept string);
```

```
hive> drop table student;
```

```
hive> use default;
```

```
hive> drop table student;
```

```
hive> use info;
```

```
hive> create table student(roll_no int, name string,dept string)
```

```
> row format delimited
```

```
> fields terminated by ',';
```

```
hive> describe student;
```

```
OK
```

roll_no	int
name	string
dept	string

For second terminal

```
[cloudera@quickstart ~]$ dir
```

```
[cloudera@quickstart ~]$ cd Desktop
```

```
[cloudera@quickstart Desktop]$ mkdir hive_data
```

```
[cloudera@quickstart Desktop]$ cd hive_data
```

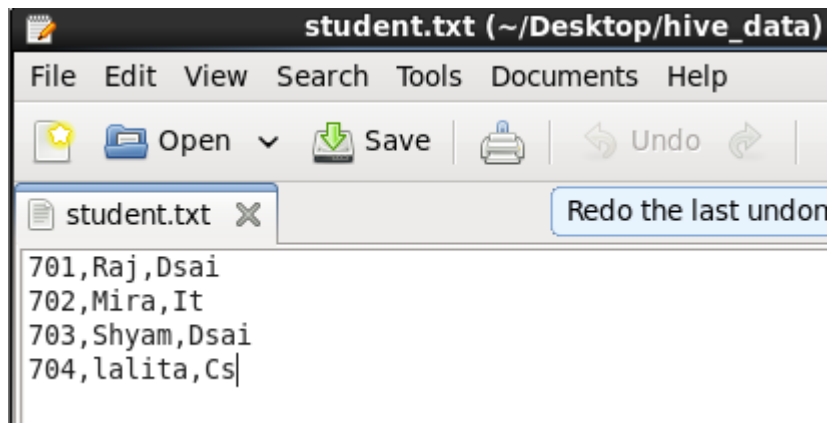
```
[cloudera@quickstart hive_data]$ dir
```

```
[cloudera@quickstart hive_data]$ touch student.txt
```

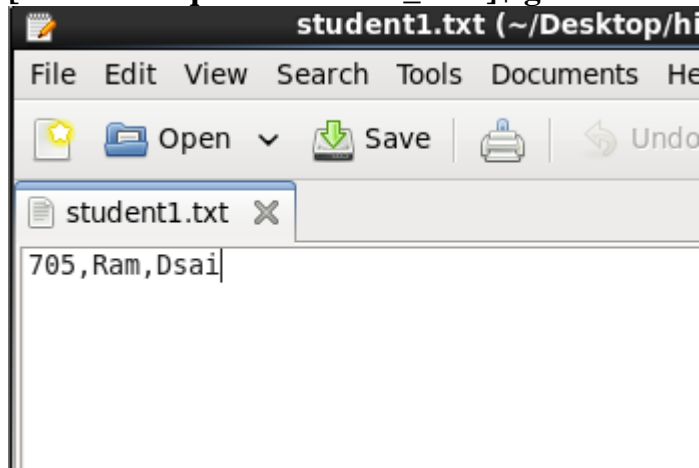
```
[cloudera@quickstart hive_data]$ dir
```

```
student.txt
```


```
[cloudera@quickstart hive_data]$ gedit student.txt
```



[cloudera@quickstart hive_data]\$ gedit student1.txt



[cloudera@quickstart hive_data]\$ gedit student2.txt



[cloudera@quickstart hive_data]\$ pwd
/home/cloudera/Desktop/hive_data....(path)

First terminal

hive> load data local inpath '/home/cloudera/Desktop/hive_data' into table student;

```
hive> load data local inpath '/home/cloudera/Desktop/hive_data' into table student;
Loading data to table info.student
Table info.student stats: [numFiles=4, numBytes=100]
OK
```

hive> select * from student;

```
OK
701    Raj    Dsai
702    Mira    It
703    Shyam  Dsai
704    lalita  Cs
701    Raj    Dsai
702    Mira    It
703    Shyam  Dsai
704    lalita  Cs
705    Ram    Dsai
Time taken: 0.302 seconds
```

hive> truncate table student;

hive> select * from student;

hive> load data local inpath '/home/cloudera/Desktop/hive_data/student.txt'
into table student;

hive> select * from student;

```
OK
701    Raj    Dsai
702    Mira    It
703    Shyam  Dsai
704    lalita  Cs
Time taken: 0.055 seconds,
```

load data local inpath '/home/cloudera/Desktop/hive_data/student.txt' into
table student;

hive> select * from student;

```
OK
701    Raj    Dsai
702    Mira    It
703    Shyam  Dsai
704    lalita  Cs
```

hive> load data local inpath '/home/cloudera/Desktop/hive_data/student2.txt'
into table student;

hive> select * from student;

```
OK
701    Raj    Dsai
702    Mira    It
703    Shyam  Dsai
704    lalita  Cs
709    Don    420
710    Hina    916
```

hive> describe extended student;

```
OK
roll_no      int
name         string
dept         string
```

hive> describe extended ext_student;

```
OK
rol_no      int
name        string
dept        string
```

tableType:EXTERNAL_TABLE).....

hive> drop table ext_student;

hive> use info;

Terminal 3

[cloudera@quickstart ~]\$ hdfs dfs -mkdir /hadoop_data

[cloudera@quickstart ~]\$ hdfs dfs -ls /

Terminal 1

create table ext_student(roll_no int, name string,dept string)

> row format delimited

> fields terminated by ','

> location '/hadoop_data';

hive> select *from ext_student;

```
OK
701    Raj    Dsai
702    Mira   It
703    Shyam  Dsai
704    lalita Cs
701    Raj    Dsai
702    Mira   It
703    Shyam  Dsai
704    lalita Cs
705    Ram    Dsai
709    Don    420
710    Hina   916
```

Practical 8

Perform Data Ingestion using Apache Sqoop tool:

- a) MySQL to HDFS (import)
- b) MySQL to HIVE (import)
- c) MySQL to HBASE (import)
- d) HDFS to MySQL (export)
- e) HIVE to MySQL (export)

```
[cloudera@quickstart ~]$ mysql -u root -pcloudera
mysql> show databases;
```

Database
information_schema
cm
firehose
hue
metastore
mysql
nav
navms
oozie
retail_db
rman
sentry

```
mysql> create database student;
```

Database
information_schema
cm
firehose
hue
metastore
mysql
nav
navms
oozie
retail_db
rman
sentry
student

```
mysql> create database rjc;
mysql> show databases;
```

Database
information_schema
cm
firehose
hue
metastore
mysql
nav
navms
oozie
retail_db
rjc
rman
sentry
student

```
mysql> use rjc;
mysql> create table student_info(id int(10),name char(20),address varchar(50));
mysql> show tables;
```

Tables_in_rjc
student_info

```
mysql> describe student_info;
```

Field	Type	Null	Key	Default	Extra
id	int(10)	YES		NULL	
name	char(20)	YES		NULL	
address	varchar(50)	YES		NULL	

```
mysql> insert into student_info values(1,'Ramesh','Gatkhpar');
mysql> insert into student_info values(2,'Suresh','Mumbai');
mysql> insert into student_info values(3,'Mohit','Chennai');
mysql> select * from student_info
```

id	name	address
1	Ramesh	Gatkhpar
2	Suresh	Mumbai
3	Mohit	Chennai

```
mysql> create table course_info(id int(10),course_name char(20),course_teacher
varchar(50))
mysql> describe course_info;
```

Field	Type	Null	Key	Default	Extra
id	int(10)	YES		NULL	
course_name	char(20)	YES		NULL	
course_teacher	varchar(50)	YES		NULL	

```
mysql> insert into course_info values(1,'python','Mujtaba sir');
mysql> insert into course_info values(2,'Data em','Neha mam');
mysql> insert into course_info values(3,'ML','rahul sir');
mysql> select * from course_info;
```

id	course_name	course_teacher
1	python	Mujtaba sir
2	Data em	Neha mam
3	ML	rahul sir

```
mysql> create table Department(id int(10),dept_name char(20),no_of_teacher
int(10));
create table Department(id int(10),dept_name char(20),no_of_teacher int(10));
mysql> select * from Department;
mysql> describe Department;
```

Field	Type	Null	Key	Default	Extra
id	int(10)	YES		NULL	
dept_name	char(20)	YES		NULL	
no_of_teacher	int(10)	YES		NULL	

```
mysql> insert into Department values(1,'data sci',6);
mysql> insert into Department values(1,'data sci',6);
mysql> insert into Department values(3,'info',12);
mysql> select * from Department;
```

id	dept_name	no_of_teacher
1	data sci	6
2	com Sci	12
3	info	12

Terminal 2

```
[cloudera@quickstart ~]$ whoami
[cloudera@quickstart ~]$ hostname
[cloudera@quickstart ~]$ sqoop list-databases --connect
jdbc:mysql://quickstart.cloudera:3306/ --password cloudera
--username root;
```

information_schema
cm
firehose
hue
info
info
metastore
mysql
nav
navms
oozie
retail_db
rman
sentry

From sql to hdfs(import)

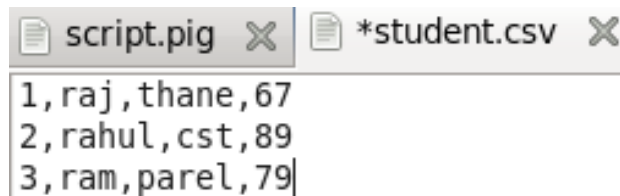
```
[cloudera@quickstart ~]$ sqoop import --connect  
jdbc:mysql://quickstart.cloudera:3306/info --password cloudera --username root --  
table stu --m 1;  
From
```

```
[cloudera@quickstart ~]$ hdfs dfs -cat /user/cloudera/student_info/part*
```


Practical 9

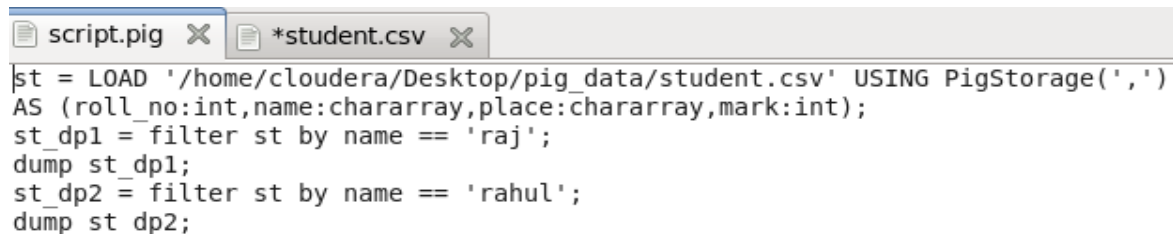
Perform data processing using pig latin.

[cloudera@quickstart pig_practical]gedit student.csv

A screenshot of a gedit window with two tabs: 'script.pig' and '*student.csv'. The 'student.csv' tab is active, showing the following content:

```
1,raj,thane,67
2,rahul,cst,89
3,ram,parel,79
```

[cloudera@quickstart pig_practical]gedit script.pig

A screenshot of a gedit window with two tabs: 'script.pig' and '*student.csv'. The 'script.pig' tab is active, showing the following Pig Latin script:

```
st = LOAD '/home/cloudera/Desktop/pig_data/student.csv' USING PigStorage(',')
AS (roll_no:int,name:chararray,place:chararray,mark:int);
st_dp1 = filter st by name == 'raj';
dump st_dp1;
st_dp2 = filter st by name == 'rahul';
dump st_dp2;
```

Write on terminal pig -x local

exec script.pig

```
2024-03-02 07:40:
ne.util.MapRedUti
(2,rahul,cst,89)
grunt> █
```

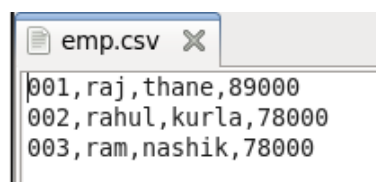
terminal 1

[cloudera@quickstart ~]\$ cd Desktop

[cloudera@quickstart Desktop]\$ cd pig_practical

[cloudera@quickstart pig_practical]\$ hdfs dfs -mkdir /big_practical

[cloudera@quickstart pig_practical]\$ gedit emp.csv

A screenshot of a gedit window with one tab: 'emp.csv'. The file contains the following content:

```
001,raj,thane,89000
002,rahul,kurla,78000
003,ram,nashik,78000
```

[cloudera@quickstart pig_practical]\$ hdfs dfs -put emp.csv /big_practical

Terminal 2 pig

```
grunt> emp = LOAD '/big_practical/emp.csv' USING PigStorage(',') AS  
(emp_id:int,emp_name:chararray,emp_loc:chararray,emp_salary:int);  
grunt> dump emp;
```

```
org.apache.pig.util.MapRedUtil  
(1,raj,thane,89000)  
(2,rahul,kurla,78000)  
(3,ram,nashik,78000)  
(,)
```

```
grunt> emp_dp = filter emp by emp_name == 'raj';  
grunt> dump emp_dp;
```

```
org.apache.pig.util.MapRedUtil  
(1,raj,thane,89000)
```

```
grunt> emp_dp2 = foreach emp generate emp_salary,emp_id;  
grunt> dump emp_dp2;
```

```
org.apache.pig.util.MapRedUtil  
(89000,1)  
(78000,2)  
(78000,3)  
(,)
```

```
grunt> emp_dp3 = order emp by emp_salary ;  
grunt> dump emp_dp3 ;
```

```
org.apache.pig.util.MapRedUtil  
(3,ram,nashik,78000)  
(2,rahul,kurla,78000)  
(1,raj,thane,89000)
```